

Unequal benefits of AI in the labor market: A gender perspective

Claudia Romano

<i>1. Introduction</i>	<i>2</i>
<i>2. Displacement and productivity: Acemoglu et al. study</i>	<i>2</i>
<i>2.1 The task-based theoretical model</i>	<i>2</i>
<i>2.2 Empirical strategy</i>	<i>3</i>
<i>2.3 Key findings</i>	<i>3</i>
<i>3. Polarization of the labor market: Felten et al. study</i>	<i>6</i>
<i>3.1 AIOI score index</i>	<i>6</i>
<i>3.2 The empirical strategy</i>	<i>7</i>
<i>3.3 Key findings</i>	<i>7</i>
<i>4. Gender inequality</i>	<i>8</i>
<i>5. Policy implications</i>	<i>9</i>

1. Introduction

AI as a new and always improving technology is still under the lens, the scientific literature is mostly newborn if we compare it with the ones about other economic problems. There is the need to create and use new measures to actually understand the extent of its impact. Although it may be still too soon to discover the long run effects or the clear macroeconomic patterns, researchers remain committed to exploring these dynamics, also in order to orient policy makers. Fundamental questions persist regarding the true beneficiaries of AI-driven progress. Specifically, while the automation of traditionally human tasks offers efficiency gains, it simultaneously poses the risk of severe labor market disruption and widening inequality. Determining exactly which occupational groups and social classes are benefiting from this technological transition is imperative.

2. Displacement and productivity: Acemoglu et al. study

2.1 The task-based theoretical model

In *Artificial intelligence and jobs: evidence from Online Vacancies* (Acemoglu et al. 2022), the authors conduct an establishment-level research on the fields of online vacancies in the US from 2010 to 2018, dividing the timeline into two blocks: 2010-2012 and 2016-2018, focusing on firms using AI to complete tasks (not on firms supplying AI services). The measure of AI exposure of a firm from its baseline (2010-2012) occupational structure is constructed starting from three different indices as proxies: Felten et al. AI Occupational Impact measure, SML index, and Webb's (2020) AI exposure score. The aim of this paper is to study the impact of AI adoption in US firms on the labor market starting from a theoretical task-based model.

It assumes that each establishment has an increasing demand function for its product with its price and wants to maximize the profit having small market power and taking other firms' prices as given.

The first proposition of the model affirms that the changes in AI activity and hiring of AI workers are both proportional to exposure to AI. As a consequence, the posting of more vacancies for in-house AI workers is dependent on the variable exposure to AI in the empirical analysis.

The second proposition affirms that for an always increasing technological progress and increasing number of tasks profitably achievable by AI, with the assumption of perfect substitutability of AI and labor for certain tasks, there is a displacement effect for workers, that reduces employment.

The displacement effect in theory could be moderated by the productivity effect, generated by the decreasing of costs due to automatized labor and consequent efficiency increasing, that leads to more profits and an increasing human labor demand in non AI-related jobs. If we look only at the theory, it is impossible to know ex ante which of the two effects will prevail.

2.2 Empirical strategy

The empirical strategy ties AI exposure, the independent variable, to the three main outcome measures, the dependent variables: AI vacancies, changes in job skill requirements of posted vacancies, and overall non AI hiring.

In order to do that, the authors estimate the following regression model:

$$\Delta Y_{e,t-t_0} = \beta AI_{e,t_0} + x'_{e,t_0} \gamma + \varepsilon_{e,t-t_0} \text{ where}$$

$\Delta Y_{e,t-t_0}$ is the change in the establishment-level outcomes between 2010-2012 and 2016-2018;

AI_{e,t_0} is the exposure of the establishment e to AI in the period 2010-2012;

$x'_{e,t_0} \gamma$ is the controls vector, like industry fixed effects, commuting zones fixed effects, size of firms and so on;

$\varepsilon_{e,t-t_0}$ is the error;

β is the crucial coefficient of the analysis that states how much the outcome changes based on the increasing in standard deviation of in AI exposure; if it is positive, the outcome has an increase when AI exposure does so; if it is negative, the outcome decreases when AI exposure increases.

2.3 Key findings

Growth of Establishment AI Vacancies, 2010–18						
	(1)	(2)	(3)	(4)	(5)	(6)
A. Felten et al. Measure of AI Exposure						
Establishment AI Exposure, 2010	15.96 (1.73)	13.82 (1.43)	9.19 (1.21)	16.53 (1.89)	9.75 (1.20)	16.87 (1.86)
Observations	1,075,474	1,075,474	954,519	1,075,474	954,518	762,672
B. Webb Measure of AI Exposure						
Establishment AI Exposure, 2010	6.59 (1.13)	5.08 (.96)	3.21 (.81)	5.91 (1.27)	.42 (.82)	1.14 (1.08)
Observations	1,159,789	1,159,789	1,021,673	1,159,789	1,021,673	824,803
C. SML Measure of AI Exposure						
Establishment AI Exposure, 2010	3.76 (1.19)	2.30 (1.04)	−2.21 (.96)	−3.04 (1.38)	1.95 (.89)	4.47 (1.34)
Observations	1,159,789	1,159,789	1,021,673	1,159,789	1,021,673	824,803
Covariates:						
Share of vacancies in sales and administration, 2010					✓	✓
Fixed effects:						
Firm size decile		✓	✓		✓	
CZ		✓	✓	✓	✓	✓
Three-digit industry			✓		✓	
Firm				✓		✓

Table 1: Acemoglu et al. (2022). Effects of AI Exposure on Establishment AI Vacancy Growth, 2010–18

The authors find a robust, statistically significant positive correlation between baseline AI exposure and subsequent growth in AI-related vacancies, both for Felten and Webb measures (with SML the phenomenon is less apparent). Establishments whose task structures enable

the use of AI technologies have substantially increased their AI-related postings, evidencing that the type of task has a main role in adoption of AI tools. This is valid also when controlling for exposure to “traditional” software.

Growth of Establishment AI Vacancies, 2010–18						
	(1)	(2)	(3)	(4)	(5)	(6)
A. Felten et al. Measure of AI Exposure						
Establishment AI Exposure, 2010	16.28 (1.74)	14.10 (1.44)	9.63 (1.23)	17.43 (1.95)	9.96 (1.24)	17.47 (1.90)
Establishment software Exposure	2.36 (.76)	2.24 (.71)	2.62 (.76)	4.83 (1.23)	.66 (.82)	1.83 (1.19)
Observations	1,059,620	1,059,620	941,046	1,059,620	941,046	1,059,620
B. Webb Measure of AI Exposure						
Establishment AI Exposure, 2010	10.64 (1.83)	7.88 (1.50)	3.85 (1.10)	6.81 (1.52)	1.50 (1.07)	2.57 (1.41)
Establishment software Exposure	-6.28 (1.51)	-4.27 (1.26)	-.96 (.96)	-1.44 (1.34)	-1.81 (1.00)	-2.54 (1.49)
Observations	1,159,789	1,159,789	1,021,673	1,159,789	1,021,673	1,159,789
C. SML Measure of AI Exposure						
Establishment AI Exposure, 2010	4.14 (1.18)	2.64 (1.04)	-1.96 (.96)	-2.42 (1.28)	1.90 (.88)	4.37 (1.28)
Establishment software Exposure	1.56 (.80)	1.44 (.77)	1.09 (.69)	2.40 (1.04)	-.90 (.78)	-.76 (1.11)
Observations	1,159,789	1,159,789	1,021,673	1,159,789	1,021,673	1,159,789
Covariates:						
Share of vacancies in sales and administration, 2010					✓	✓
Fixed effects:						
Firm size decile		✓	✓		✓	
CZ		✓		✓	✓	✓
Three-digit industry			✓		✓	
Firm				✓		✓

Table 2: Acemoglu et al. (2022). Effects of AI Exposure on Establishment AI Vacancy Growth, Controlling for Software Exposure

A higher AI exposure is strongly associated with the disappearance of skills previously listed in job postings. For a 1 standard deviation increase in exposure, there is a 0.83 absolute decline in the frequency of previously demanded skills.

Establishment Negative Skill Change, 2010–18						
	(1)	(2)	(3)	(4)	(5)	(6)
A. Felten et al. Measure of AI Exposure						
Establishment AI Exposure, 2010	.83 (.09)	.83 (.09)	.97 (.07)	.50 (.05)	1.00 (.07)	.54 (.05)
Observations	339,282	339,282	322,901	339,282	322,901	339,282
B. Webb Measure of AI Exposure						
Establishment AI Exposure, 2010	.62 (.11)	.60 (.11)	.45 (.06)	.20 (.04)	.68 (.11)	.34 (.04)
Observations	353,107	353,107	335,589	353,107	335,589	353,107
C. SML Measure of AI Exposure						
Establishment AI Exposure, 2010	.53 (.08)	.52 (.07)	.32 (.07)	.26 (.04)	.46 (.09)	.36 (.04)
Observations	353,107	353,107	335,589	353,107	335,589	353,107
Covariates:						
Share of vacancies in sales and administration, 2010					✓	✓
Fixed effects:						
Firm size decile		✓	✓		✓	
CZ		✓	✓	✓	✓	✓
Three-digit industry			✓		✓	
Firm				✓		✓

Table 3: Acemoglu et al. (2022). Effects of AI Exposure on Establishment Negative Skill Change, 2010–18

Exposed establishments have a 0.95 absolute increase in the frequency of new or emerging skills.

	Establishment Positive Skill Change, 2010–18					
	(1)	(2)	(3)	(4)	(5)	(6)
A. Felten et al. Measure of AI Exposure						
Establishment AI Exposure, 2010	.95 (.08)	.94 (.09)	.58 (.09)	.02 (.04)	.62 (.09)	.05 (.04)
Observations	339,282	339,282	322,901	339,282	322,901	339,282
B. Webb Measure of AI Exposure						
Establishment AI Exposure, 2010	.69 (.09)	.66 (.09)	.26 (.08)	–.01 (.03)	.43 (.08)	.13 (.04)
Observations	353,107	353,107	335,589	353,107	335,589	353,107
C. SML Measure of AI Exposure						
Establishment AI Exposure, 2010	.62 (.09)	.59 (.09)	.19 (.09)	.10 (.04)	.26 (.09)	.03 (.04)
Observations	353,107	353,107	335,589	353,107	335,589	353,107
Covariates:						
Share of vacancies in sales and administration					✓	✓
Fixed effects:						
Firm size decile		✓	✓		✓	
CZ		✓	✓	✓	✓	✓
Three-digit industry			✓		✓	
Firm				✓		✓

Table 4: Acemoglu et al. (2022). Effects of AI Exposure on Establishment Positive Skill Change, 2010–18

The results in Table 5 reveal a consistent negative association between AI exposure and non-AI vacancy growth. Specifically, a 1 standard deviation increase in exposure correlates with an approximate 14% decline in non-AI vacancies. When these vacancy flows are converted into employment stocks, the estimates suggest a 7.2% reduction in non-AI employment.

	Growth of Establishment Non-AI Vacancies, 2010–18							
	Full Sample						Establishments Posting in 2018	
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
A. Felten et al. Measure of AI Exposure								
Establishment AI Exposure, 2010	–13.80 (4.22)	–16.36 (4.11)	–11.90 (4.08)	–4.81 (1.44)	–12.42 (4.01)	–4.04 (1.47)	–8.38 (3.46)	–3.56 (1.86)
Observations	1,075,474	1,075,474	954,519	1,075,474	954,519	1,075,474	324,901	341,525
B. Webb Measure of AI Exposure								
Establishment AI Exposure, 2010	–17.24 (3.72)	–18.21 (3.63)	–6.73 (3.01)	–2.22 (.93)	–8.30 (3.70)	1.51 (.98)	–4.70 (2.66)	–1.44 (1.36)
Observations	1,159,789	1,159,789	1,021,673	1,159,789	1,021,673	1,159,789	337,758	355,529

Table 5: Acemoglu et al. (2022). Effects of AI Exposure on Establishment Non-AI Vacancy Growth, 2010–18

These negative coefficients imply that the productivity effect does not make up for the displacement effect. Table 7 validates this by showing that these negative hiring effects are absent in the 2010–2014 period but become highly significant during the 2014–2018 AI takeoff, aligning the results with the timing of technological advancement.

Table 8 zooms out to look at the big picture, testing whether the changes happening inside individual companies are actually affecting entire industries or broad career fields.

	Industry by CZ Employment Growth (CBP)			Occupation Employment Growth (OES)			Occupation Wage Growth (OES)		
	2003–7 (1)	2007–10 (2)	2010–16 (3)	2004–7 (4)	2007–10 (5)	2010–18 (6)	2004–7 (7)	2007–10 (8)	2010–18 (9)
A. Felten et al. Measure of AI Exposure									
Market AI exposure, 2010	.03 (.17)	.10 (.20)	-.05 (.08)	.34 (.34)	.86 (.32)	.51 (.35)	-.00 (.17)	.02 (.20)	-.17 (.06)
Observations	10,937	10,926	10,929	736	700	680	680	648	629
B. Webb Measure of AI Exposure									
Market AI exposure, 2010	.10 (.15)	.18 (.17)	.11 (.09)	.00 (.17)	.11 (.21)	-.17 (.29)	.11 (.08)	-.05 (.10)	-.02 (.04)
Observations	10,981	10,968	10,968	713	704	717	660	653	663

Table 8: Acemoglu et al. (2022). Effects of AI Exposure on Market Employment and Wage Growth

Unlike the strong results found at the company level, the data shows that overall job growth and average wages have barely moved; the statistical numbers are essentially zero and show no clear pattern. While the results are significant at the establishment-level the impact is still too small relative to the scale of the US labor market to trace back an employment pattern.

3. Polarization of the labor market: Felten et al. study

3.1 AIOI score index

The AIOI score index creation is explained in *The Occupational Impact of Artificial Intelligence: Labor, Skills, and Polarization* (Felten et al. 2019). The focus is the link between AI and wages, employment and labor market polarization. It's created linking the Electronic Frontier Foundation, that traces AI rates of progress in 9 specific applications from 2010 to 2015 (as shown in table 1) to the O*NET dataset, which categorizes professions/occupations describing them through 52 specific human abilities, to which one is assigned a score of "importance" from 1 to 5 and of "prevalence" from 1 to 5 in every occupation. To connect the AI applications to human abilities, the authors used a crowdsourced sample of approximately 1,800 "gig workers" from Amazon Mechanical Turk (mTurk). For each of the 9 AI applications, 200 respondents answered a simple "Yes" or "No" regarding whether they believed the AI application was related to or could be used for each of the 52 O*NET abilities. These binary responses were averaged to construct a matrix measuring the relatedness between each AI application and each ability. Then they multiplied the calculated rate of progress for each AI application by the application-ability relatedness measure from the matrix and summed this product across all 9 AI applications to get the total AI impact for that specific ability. Next, they evaluated the impact on specific jobs using O*NET occupational definitions. They weighted the ability-level AI impact by multiplying it by the scaled prevalence and importance scores of that specific ability within a given occupation. They aggregated the total impact on an occupation by summing these weighted ability-level AI impacts across all the abilities required for that job. They scaled this aggregated impact by dividing it by the weighted sum of the prevalence and importance of all abilities used in the occupation.

Finally, the O*NET data organizes occupations at a very granular 8-digit Standard Occupational Classification (SOC) level (832 occupations). To easily use the AIOI scores with other major data sources, the researchers collapsed them into 742 occupations at the

standard 6-digit SOC level by calculating the mean AIOI score across the grouped occupations.

The first comparison is of two professions that share some similarities in terms of the ability set (manual dexterity, finger dexterity, and arm-hand steadiness) required yet have highly disparate AIOI scores: surgeons and meat slaughterers. The firsts are more affected by AI and this difference seems to stem from the cognitive abilities related to problem solving required by each occupation. What this example is suggesting is that, given the limited impact of AI on physical abilities, the impact of AI on occupations largely stems from cognitive abilities like problem sensitivity, deductive and inductive reasoning, flexibility of closure, and information ordering.

3.2 The empirical strategy

Now, before the analysis at the occupational-level, the point is to understand whether AI is having a role of labor substitution or complementarity. This is done using employment and wage at state level data from BLS from 2010 to 2016 for each occupation.

The regression models take the following form:

$$Y_{is} = \beta_0 + \beta_1 AIOI\ score_i + \beta_2 X_i + \gamma_s + \varepsilon_{is}$$

where:

Y: employment growth or wage growth;

X: vector of controls;

γ_s : fixed effect for the state;

i: occupation;

s: state;

β_1 : coefficient of interest;

errors clustered at the state level.

Completely new occupations that are created thanks to innovation will not be captured. For each regression, they control for employment or wage growth from 2005 to 2010, depending on the dependent variable to account for pre-existing trends. Then the authors construct two additional variables: Degree of Automation and Importance of Programming. In some of our regressions, then include the variable State Population Change.

3.3 Key findings

Table 7 of the paper (here table 9) shows the outcomes of the relationships of interest.

At first, we can see that AIOI does not have a statistically significant and meaningful relationship with employment growth, while an increase of 1 in the standard deviation of AIOI is related with a 3.5 percent increase in wage growth (scaling by the sample mean of the median wage growth). Based on occupation-specific characteristics, they classify occupations into terciles by the proportion of software skills required.

	Employment Growth (2010-2016)				Wage Growth (2010-2016)			
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
AIOI Score	-0.0254 (0.1409)	-0.0292 (0.1432)	-0.0463 (0.1458)	0.1231 (0.1512)	0.1367*** (0.0367)	0.1345*** (0.0368)	0.1311*** (0.0367)	0.1160*** (0.0402)
Employment Growth (2005-2010)	-0.1443*** (0.0122)	-0.1472*** (0.0122)	-0.1468*** (0.0124)	-0.1485*** (0.0125)				
Wage Growth (2005-2010)					-0.3349*** (0.0129)	-0.3364*** (0.0125)	-0.3458*** (0.0126)	-0.3505*** (0.0129)
Degree of Automation				-0.0035*** (0.0003)				-0.0000 (0.0001)
Importance of Programming				0.0021*** (0.0004)				0.0003*** (0.0001)
State Population Change (2010-2016)		1.2400*** (0.2185)				0.1944 (0.1495)		
Constant	0.1840* (0.0962)	0.1359 (0.0968)	0.1983** (0.0980)	0.1160 (0.0990)	0.0772*** (0.0261)	0.0709*** (0.0257)	0.0827*** (0.0247)	0.0843*** (0.0256)
State FEs	No	No	Yes	Yes	No	No	Yes	Yes
Observations	25,485	25,485	25,485	25,143	27,899	27,899	27,899	27,497
R-squared	0.0191	0.0233	0.0288	0.0342	0.0003	0.0010	0.0177	0.0196

Clustered standard errors in parentheses
*** p<0.01, ** p<0.05, * p<0.1

Table 9: Felten et al. (2019). Relationship between AIOI and Employment and Wage Growth (2010 – 2016)

Completing the analysis of the other outcomes in the study, the positive relationship between AI and wage growth is largely driven by occupations involving a high level of software skills. Key results show that AI may have a particularly positive impact on high-income occupations. Differently than before, an increase of 1 in the standard deviation of AIOI score is associated with an increase in employment growth of 3.94 percentage points. At least in the short run analyzed, AI does not seem to be a factor of substitution to labor, but of complementarity. However, it could be widening the income inequality and exacerbating labor polarization, while its benefits are concentrated in favor of high-income occupations. Both papers support the theory that AI adoption and gainings are unequally distributed and the productivity premium does not represent a benefit for everyone. Even when complementary to workers, it is complementary and serves only few workers, only few sectors, only few categories. Another example of this unequal effect is the difference in impact across women and men.

4. Gender inequality

From the OECD policy brief “*Algorithm and Eve: How AI will impact women at work*” on 22 countries, we can clearly see two important statements on this issue, moving from the US point of view of the previous papers analyzed.

The first one is that women, besides having the same AI exposure at work as men do, are underrepresented in the occupations with the very highest exposure: science related jobs that require tertiary education. That’s not news, women have always suffered the divide in participation in STEM jobs, due to stereotyping, lack of role models, discrimination in recruiting processes and adverse working conditions. Although women’s employment growth is greater than men’s, it’s more to be attributed to declining gender segregation than to AI. It is serious per se, but it also implies that women are underrepresented in decision-making processes about AI, increasing the risk of total or partial omissions of their voices and their experiences and of biasedness of AI tools. A technology made and trained by excluding the same sociodemographic groups that have always been penalized will clearly perpetuate the same gender stereotypes that have always permeated the labor market. Underrepresentation of any vulnerable group in decision-making roles developing and implementing AI increases

the risk that the experiences and voices of these groups are omitted from the process (Lane, 2024).

The second statement is that women are overrepresented in clerical jobs, occupations that are at higher risk of automation.

Furthermore, women are penalized even in the very access to employment in workplaces where automated screening procedures are used.

With up to 98.4% of Fortune 500 companies incorporating AI into their hiring processes, employment discrimination can propagate at an unprecedented scale. A simulation of Large Language Model (LLM) resume screening by Wilson and Caliskan (Brookings Institution, 2025) demonstrated that systems favored men's names 51.9% of the time, while preferring women's names in only 11.1% of cases.

The severity of this bias requires an intersectional approach to accurately quantify societal harm, since the combination of multiple identities leads to distinct outcomes that differ from experiencing any single identity independently.

At the intersection of race and gender, LLM screening selected resumes with Black men's names 0% of the time when compared directly to those with white men's names. Simple debiasing techniques, like removing explicit names or gender references, fail because models infer social identities through implicit cues like attendance at an all-women's college.

No matter how objective or impartial this technology can seem, it needs regulation in order to avoid perpetuating or creating biases. Of course, this won't erase gender inequalities once and for all. For example, if an AI is trained to select candidates based on experience years and a woman has been involuntarily unemployed for years due to other gender-based reasons (like incompatibility with the maternity role), the real problem stems from an unequal culture and world, not from an automated process.

The topic is multifaceted and it requires critical thinking and training of those who are responsible for recruitment processes.

Following this reasoning, gender penalties in historically male-dominated fields reflect systemic inequalities that have existed long before the current capabilities of artificial intelligence were even imaginable.

In light of these considerations, nowadays technology is one of the main drivers of progress in the way the system intends it and it should be included in the way we think about intersectionality in decision-making processes.

Vulnerable groups should be taken into account in decision making in the near future: non-native speakers in a country, workers with lower degree of education, people with disabilities, people in fragile regions of the world.

5. Policy implications

Policymakers who try to bring to light AI's potential need to be aware that AI is not a neutral technology.

The implementation of AI's tool should raise some concerns about who is actually benefiting from it and who, instead, is being left behind.

A first regulation step in Europe has been taken with the AI act (Regulation 2024/1689) in order to harmonise the rules and address the concerns about disadvantageous outcomes (especially in the labor market), health and medical screening, security, human and civil rights, putting at the center of the narrative human decisions and oversight and control of AI tools (human-in-the-loop) to avoid discrimination bias carried out by algorithms in processes of employment, workers' management and access to self-employment. In particular, those regarding recruiting, hiring, promotions, evaluations, task allocation and automated termination are classified as high-risk AI systems. The goal is to ensure privacy of workers, transparency, a fairer and without bias recruiting and other protections for employees. This framework is still not completely in force, so it's difficult to say if it will be effective in the long run and which effects will prevail in the labor market, especially in terms of gender equality. Criticisms about it concern the lack of impact assessment measures of gender inequality to address the systematic and specific discrimination of women. An example could be the Gender Impact Assessment and the use of gender-sensitive indicators to measure impacts in areas such as labour market participation, examining intersectionality and applying a feminist lens. The act could not meet the expectations, because ultimately AI systems depend on data, causing data bias resulting from unrepresentative, incomplete or historically biased datasets and perpetuating existing inequalities (Karagianni, 2025). According to Lane, the gender lens must concern upskilling and reskilling to equip workers with the knowledge required to use AI effectively, but most importantly closing the gap between women and men in tech. Inclusion of diverse and marginalized groups is the crucial point, starting from the education of girls: mapping gender inequality in STEM, incorporating female models and giving incentives are the three main pillars. This might not fully address the issue and it is necessary to integrate strong financial support and unemployment benefits to protect vulnerable workers during lengthy retraining periods, making active and passive labor market policies coordinate to ensure their full efficiency. In conclusion, ensuring that the benefits of artificial intelligence are equitably distributed, rather than monopolized by a privileged few, requires robust regulatory and monitoring frameworks across all levels of governance. This multi-level approach is particularly critical in decentralized contexts like Italy, where regional heterogeneity and devolved competencies in crucial sectors such as healthcare (devolved to regions) and education (in a context with strong autonomy of universities) play a fundamental role.

References

(2024) Algorithm and Eve | OECD. Available at: https://www.oecd.org/en/publications/2024/12/algorithm-and-eve_0e889c45.html (Accessed: 22 May 2026).

Acemoglu, D. et al. (2022) ‘Artificial Intelligence and jobs: Evidence from online vacancies’, *Journal of Labor Economics*, 40(S1). doi:10.1086/718327.

Felten, E.W., Raj, M. and Seamans, R. (1970) The occupational impact of Artificial Intelligence: Labor, skills, and polarization, *For Work / Against Work*. Available at: <https://onwork.edu.au/bibitem/2019-Felten,Edward+W-Raj,Manav-et-al-The+Occupational+Impact+of+Artificial+Intelligence+Labor,Skills,and+Polarization/> (Accessed: 20 May 2026).

Karagianni, A. (2025) The EU Artificial Intelligence Act through a gender lens, *Publikationen der Stiftung / The EU Artificial Intelligence Act through a gender lens*. Available at: <https://collections.fes.de/publikationen/ident/fes/21887> (Accessed: 20 May 2026).

Kyra Wilson, A.C. et al. (2025) Gender, race, and intersectional bias in AI resume screening via language model retrieval, *Brookings*. Available at: <https://www.brookings.edu/articles/gender-race-and-intersectional-bias-in-ai-resume-screening-via-language-model-retrieval/> (Accessed: 20 May 2026).

Lane, M. (2024) ‘Who will be the workers most affected by AI?’, *OECD Artificial Intelligence Papers* [Preprint]. doi:10.1787/14dc6f89-en.